

The OS section tests your understanding of how a computer manages hardware and software resources. The syllabus focuses on the following key areas:

A. Basics & System Structure

- **Definition:** An OS acts as an interface between the user and computer hardware, managing resources like the CPU and memory.
- **Types of OS:**
 - **Multiprogramming:** Maximizes CPU utilization by keeping multiple jobs in memory. When one waits for I/O, the CPU switches to another.
 - **Time-Sharing (Multitasking):** Uses "time slices" (quanta) to switch rapidly between tasks, giving the illusion of simultaneous execution.
 - **Real-Time OS (RTOS):** Used for tasks with strict deadlines (e.g., medical devices).
- **Kernel vs. User Mode:** The **Kernel** is the core component.
 - In **User Mode**, the mode bit is **1** (unprivileged operations).
 - In **Kernel Mode** (also called System or Supervisor mode), the mode bit is **0** (privileged operations).

B. Process Management

- **Process vs. Thread:** A process is a program in execution (heavyweight), while a thread is a lightweight unit of CPU execution *within* a process.
- **Schedulers:**
 - **Long-term Scheduler:** Controls the degree of multiprogramming, deciding which processes enter the ready queue (New \$to\$ Ready state transition).
 - **Short-term Scheduler:** Selects which process gets the CPU next (Ready \$to\$ Running state transition).
 - **Dispatcher:** The module that actually gives control of the CPU to the process selected by the short-term scheduler.
- **System Calls:** The `fork()` system call is used to create new processes.
(High-Frequency Question Topic)

C. CPU Scheduling Algorithms

This is a critical topic for numerical and conceptual questions.

- **FCFS (First-Come, First-Serve):** Simple, non-preemptive. Suffers from the **Convoy Effect**, where a long process blocks shorter ones behind it.
- **SJF (Shortest Job First):** Selects the process with the smallest execution time. It provides the **minimum average waiting time** but can lead to **Starvation** for long processes.
- **Round Robin (RR):** Preemptive scheduling using a fixed time quantum. If the time quantum is too large, RR behaves exactly like FCFS.

D. Memory Management

- **Logical vs. Physical Address:** The CPU generates a **logical address**; the **Memory Management Unit (MMU)** maps it to a **physical address**.
- **Fragmentation:**

- **Internal Fragmentation:** Wasted space *inside* a fixed-sized memory partition.
- **External Fragmentation:** Wasted space *outside* partitions (total free space exists but is not contiguous).
- **Paging:** Solves external fragmentation by dividing memory into fixed-size "frames" (physical memory) and logical memory into "pages."
- **Thrashing:** A state where the system spends more time swapping pages in and out (page faults) than executing tasks, often caused by insufficient RAM.

E. Deadlocks

A deadlock occurs when a set of processes are blocked because each is holding a resource and waiting for another resource held by someone else.

- **4 Necessary Conditions (MUST MEMORIZE):**
 1. Mutual Exclusion
 2. Hold and Wait
 3. No Preemption
 4. Circular Wait
- **Banker's Algorithm:** Used for **Deadlock Avoidance** to ensure the system remains in a "safe state."
- **Victim:** A process terminated to recover from a deadlock is called a "victim."

2. Important Examples with Explanations

Example 1: The fork() System Call

Concept: `fork()` creates a child process.

Formula: If `fork()` is called n times sequentially in a program, the **total number of processes** (parent + children) created is 2^n . The number of **child processes** is $2^n - 1$.

Scenario: A code segment executes `fork()` 3 times sequentially.

- **Calculation:** $n = 3$. Total processes = $2^3 = 8$. Total *child* processes = $8 - 1 = 7$.
- **Explanation:** The first fork creates 1 child (Total: 2). The second fork is executed by both the parent and the first child, creating 2 new children (Total: 4). The third fork is executed by all 4 processes, creating 4 new children (Total: 8).

Example 2: Calculating Turnaround Time (TAT)

Context: CPU Scheduling

Formula: $\text{Turnaround Time (TAT)} = \text{Completion Time (CT)} - \text{Arrival Time (AT)}$

Scenario: A process arrives at time $t=0$ and finishes execution at time $t=10$.

- **Calculation:** $TAT = 10 - 0 = 10\$$.
- **Explanation:** TAT is the total time a process spends in the system, from the moment it enters until it completes execution.

3. Key Concepts and High-Frequency Questions (Tips & Tricks)

- **Starvation Candidates:** SJF (Shortest Job First) and Priority Scheduling are algorithms that can cause starvation (a process waits indefinitely). FCFS and Round Robin generally prevent starvation.
- **Belady's Anomaly:** This is an exception where increasing the number of memory frames can actually *increase* the number of page faults. It is only observed in the FIFO (First-In-First-Out) page replacement algorithm.
- **Fastest Memory Allocation:** First Fit is typically the fastest memory allocation algorithm because it stops searching after finding the first available hole large enough.
- **Hardware vs. Software:** The MMU (Memory Management Unit) is a Hardware component. Loaders and Compilers are classified as System Programs.
- **Linux Kernel:** The Linux kernel file is often referred to as `vmlinuz`. Linux is an Open Source operating system.
- **Scheduling Criteria Goals:** A good scheduler aims to Maximize CPU Utilization and Throughput, but Minimize Waiting Time, Turnaround Time, and Response Time.

4. Recommended Reference Books

According to the C-CAT syllabus, the standard text for these concepts is:

- *Operating System Principles* by Silberschatz, Galvin, and Gagne (often called the "Dinosaur book").

Analogy to solidify understanding

Think of the Operating System as a **Restaurant Manager**:

- The **CPU** is the Chef.
- The **Processes** are the Orders.
- **RAM** is the Counter space for ingredients and finished dishes.
- The **Scheduler** is the Manager deciding which order the Chef cooks next.
 - **FCFS:** The Manager gives the Chef orders strictly in the order customers arrived (Convoy Effect: a large order blocks everyone else).
 - **SJF:** The Manager makes the Chef cook the quickest dishes first (Starvation: the person who ordered a 5-course meal waits forever).
 - **Round Robin:** The Manager lets the Chef work on an order for 2 minutes, then forces them to switch to the next order, ensuring fairness and attention to everyone.

Process Control Block (PCB)

The PCB is the data structure maintained by the Operating System (OS) for every process. It acts as a repository of all the information needed to manage a process, effectively serving as the process's **identity card**.

Key Fields Contained in a PCB	Explanation
Process State	The current state of the process (e.g., New, Ready, Running, Waiting, Halted).
Program Counter	Indicates the address of the next instruction to be executed for this process.
CPU Registers	The contents of all CPU registers when the process was interrupted, allowing it to be correctly restored.
CPU Scheduling Information	Priority of the process, pointers to the scheduling queues, and other scheduling parameters.
Memory Management Information	Includes base and limit registers, or page tables/segment tables information.
I/O Status Information	A list of I/O devices allocated to the process, and a list of open files.

Example (Context Switching):

When the OS switches the CPU from Process A (which is currently **Running**) to Process B, the following happens:

1. The OS **saves** the state of Process A (including the program counter and register values) into Process A's PCB.
2. The OS **loads** the state of Process B from Process B's PCB into the CPU's registers, setting the program counter to the next instruction for B.
3. Process B then resumes execution.

Without the PCB, the OS would not be able to pause and resume a process correctly, making **multitasking (time-sharing)** impossible.

File Control Block (FCB)

The FCB is a data structure maintained by the File System for every file. It contains the essential metadata about a file that the OS needs to manage it. In many operating systems, this is stored on disk and loaded into memory when the file is accessed.

Key Fields Contained in an FCB	Explanation
File Permissions	Access rights (e.g., Read, Write, Execute) for the owner, group, and others.
File Dates	Creation time, last modified time, and last access time.
File Size	The size of the file in bytes or blocks.
Location of File Data	A pointer to the starting disk block or a list of blocks that hold the actual file data.
File Owner and Group ID	Identifiers for the file's owner and the group it belongs to.

Example (File Access):

When a user attempts to open a file:

1. The OS finds the FCB for that file.
2. It checks the **File Permissions** field in the FCB to determine if the user has the necessary access rights (e.g., read permission).
3. If access is granted, the OS uses the **Location of File Data** field to retrieve the actual contents of the file from the disk.

Inter-Process Communication (IPC)

IPC refers to mechanisms provided by the operating system that allow independent processes to exchange data and synchronize their activities. This is crucial in a multi-process environment where processes often need to share information or coordinate tasks.

IPC Mechanism	Explanation and Example
Pipe	A simple, unidirectional or bidirectional data stream, typically used for communication between a parent and child process.
	Example: In a command line, the pipe symbol `
Socket	A communication endpoint that enables processes on different machines (or the same machine) to communicate over a

	network.
	Example: A Web Browser process communicates with a Web Server process over the internet using a pair of sockets. The server listens on a well-known port (like 80 for HTTP), and the browser connects to it via a client socket.
Semaphore	A synchronization tool (a simple integer variable) used to control access to a common resource by multiple processes. It solves the Critical Section Problem .
	Example: If two processes (P1 and P2) try to update a shared bank account balance simultaneously, a semaphore ensures that only one process can execute the update code (the critical section) at a time, preventing data corruption.

- **Focus on "Starvation":** Remember that **SJF** and **Priority Scheduling** cause starvation, while **Round Robin** and **FCFS** generally do not 10.

Starvation is a critical issue in operating systems where a process is perpetually denied access to a resource (in this case, the CPU) it needs to complete its task, even though the resource may be available at various times. The process is effectively kept waiting indefinitely. Algorithms that **Cause Starvation**

These algorithms make scheduling decisions based on a factor (like priority or time) that can consistently favor other processes.

- **SJF (Shortest Job First):**
 - **Mechanism:** SJF selects the process with the smallest *next* CPU burst time.
 - **Why it Causes Starvation:** In a system with a continuous stream of small, short jobs, the long-running process will repeatedly be preempted or bypassed. A new, shorter process will always appear and get to run before the older, longer process, leading to the long process waiting forever.
- **Priority Scheduling:**
 - **Mechanism:** Processes are executed based on a static priority level (e.g., 0 is the highest, 5 is the lowest).
 - **Why it Causes Starvation:** A low-priority process may never get a chance to run if there is always a steady supply of high-priority processes ready in the queue.
 - **Solution: Aging**

- The common solution to starvation in priority scheduling is **Aging**, where the priority of a waiting process is gradually increased over time. Eventually, its priority will become high enough that it will be selected for execution, preventing infinite waiting.

Algorithms that **Prevent** Starvation

These algorithms incorporate a mechanism to ensure fairness, preventing any single process from being permanently ignored.

- **FCFS (First-Come, First-Serve):**
 - **Mechanism:** Processes are executed in the exact order they arrive in the ready queue.
 - **Why it Prevents Starvation:** Since a process is guaranteed to be scheduled once all processes that arrived before it are completed, it cannot wait indefinitely. *The issue with FCFS is the **Convo Effect**, not starvation.*
 - **Round Robin (RR):**
 - **Mechanism:** Each process is given a small, fixed amount of time (a time quantum) to execute. If it doesn't finish, it's preempted and moved to the back of the ready queue.
 - **Why it Prevents Starvation:** RR is inherently fair. Every process in the queue receives a turn within a predictable time limit, ensuring that no process is ever left out of the rotation.
 -
 - **Hardware vs. Software:** Questions often ask if components like the **MMU** or **Drivers** are hardware or software. (MMU = Hardware; Driver = Software)
The context in your document provides a clear distinction between these two key operating system components:
 - **MMU (Memory Management Unit): Hardware**
 - **Explanation:** The MMU is a physical circuit on the computer chip (often integrated into the CPU) responsible for translating the **logical addresses** generated by the CPU into **physical addresses** in main memory (RAM).
 - **Drivers (Device Drivers): Software**
 - **Explanation:** A driver is a type of **System Program** that enables the operating system to interact with a specific hardware device (like a printer, graphics card, or mouse). It provides the software interface for the hardware.
 - **Process vs. Thread:** A process is an execution instance (heavyweight), while a thread is a lightweight unit of execution within a process
- Process:** A program in execution. It is **heavyweight** because it is an independent execution environment with its own dedicated memory space and resources, making creation and context switching slower.
- **Thread:** A lightweight unit of CPU execution *within* a process. It is **lightweight** because it shares the parent process's memory space, which allows for faster

creation and context switching.

-